Note for Reinforcement Learning 2nd Edition

July 27, 2018

Finite Markov Decision Process

MDP framework consists of an environment and agent. For $t = 0, 1, 2...t$, agent receives observed state $S_t \in \mathcal{S}$ based on which the agent perform an action $A_t \in \mathcal{A}$. The dynamic of the environemtn then returns a reward $R_{t+1} \in \mathcal{R}$. This forms a trajectory $S_0, A_0, R_1, S_1, A_1, R_2, \ldots$. The dynamic for MDP is defined to be

$$p(s', r \mid s, a) = Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\}$$

Several commonly use quantities:
state-transition probability

$$p(s' \mid s, a) = \sum_r p(s', r \mid s, a)$$

expected reward for state-action pair

$$\begin{aligned} r(s, a) &= E[R_t \mid S_{t-1} = s, A_{t-1} = a] \\ &= \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r \mid s, a) \end{aligned}$$

expected reward for state-action-next-state triples

$$\begin{aligned} r(s, a, s') &= E[R_t \mid S_{t-1} = s, A_{t-1} = a, S_t = s'] \\ &= \sum_{r \in \mathcal{R}} r p(r \mid s, a, s') \\ &= \sum_{r \in \mathcal{R}} r \frac{p(s', r \mid s, a)}{p(s' \mid s, a, s')} \end{aligned}$$

For episodic tasks, we have nonterminal states $\mathcal{S}$ and terminal states $\mathcal{S}^+$. Time of termination T is a random variable between episodes.

Returns at time step t is

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \cdots \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned}$$

All reinforcement learning algorithm involve estimating 3 quantities: value function, state-action function.

A policy is a probability distribution on $\mathcal{A}$ given s denoted as $\pi(a \mid s)$.

A value function of a state s under a policy π is the expected return when starting in s and following π thereafter.

$$\begin{aligned} v_\pi(s) &= E_\pi[G_t \mid S_t = s] \\ &= E_\pi\left[\sum_{i=0}^{\infty} \gamma^i R_{t+i+1} \middle| S_t = s\right], \quad \text{(for all } s \in \mathcal{S}) \\ &= E_\pi\left[R_{t+1} + \gamma \sum_{i=0}^{\infty} \gamma^i R_{(t+1)+i+1} \middle| S_t = s\right] \\ &= E_\pi[R_{t+1} \mid S_t = s] + \gamma E_\pi[G_{t+1} \mid S_t = s] \\ &= \sum_{a,r,s'} r p(s', r \mid s, a)\pi(a \mid s), \quad \text{(Law of Total Expectation)} \\ &\quad + \gamma \sum_{a,s'} E_\pi[G_{t+1} \mid S_{t+1} = s'] p(s' \mid s, a)\pi(a \mid s) \\ &= \sum_{a,r,s'} r p(r \mid s, a)\pi(a \mid s) + \gamma \sum_{a,r,s'} v_\pi(s') p(s', r \mid s, a)\pi(a \mid s) \\ &= \sum_a \pi(a \mid s) \sum_{r,s'} p(s', r \mid s, a)\left[r + \gamma \sum_{s'} v_\pi(s')\right] \end{aligned}$$

The state-action function is the expected return of taking an action a at state s before following the policy thereafter.

$$\begin{aligned} q(s, a) &= E_\pi[G_t \mid S_t = s, A_t = a] \\ &= E_\pi\left[\sum_{i=0}^{\infty} \gamma^i R_{t+i+1} \middle| S_t = s, A_t = a\right] \\ &= E_\pi[R_{t+1} \mid S_t = s, A_t = a] + \gamma E_\pi[G_{t+1} \mid S_t = s, A_t = a] \\ &= \sum_{s',r} r p(s', r \mid s, a) + \gamma \sum_{s',a'} E_\pi[G_{t+1} \mid s', a'] p(s', a' \mid s, a) \\ &= \sum_{s',r} r p(s', r \mid s, a) + \gamma \sum_{s',a'} q(s', a') p(a' \mid s', s, a) p(s' \mid s, a) \\ &= \sum_{s',r} r p(s', r \mid s, a) + \gamma \sum_{s',a',r} q(s', a')\pi(a' \mid s') p(s', r \mid s, a) \\ &= \sum_{s',r} p(s', r \mid s, a)\left[r + \gamma \sum_{a'} q(s', a')\pi(a' \mid s')\right] \end{aligned}$$