

Note for Reinforcement Learning 2nd Edition

July 31, 2018

Finite Markov Decision Process

MDP framework consists of an environment and agent. For $t = 0, 1, 2, \dots, t$, agent receives observed state $S_t \in \mathcal{S}$ based on which the agent perform an action $A_t \in \mathcal{A}$. The dynamic of the environment then returns a reward $R_{t+1} \in \mathcal{R}$. This forms a trajectory $S_0, A_0, R_1, S_1, A_1, R_2, \dots$. The dynamic for MDP is defined to be

$$p(s', r \mid s, a) = \Pr\{S_t = s', R_t = r \mid S_{t-1} = s, A_{t-1} = a\}$$

Several commonly use quantities:
state-transition probability

$$p(s' \mid s, a) = \sum_r p(s', r \mid s, a)$$

expected reward for state-action pair

$$\begin{aligned} r(s, a) &= E[R_t \mid S_{t-1} = s, A_{t-1} = a] \\ &= \sum_{r \in \mathcal{R}} r \sum_{s' \in \mathcal{S}} p(s', r \mid s, a) \end{aligned}$$

expected reward for state-action-next-state triples

$$\begin{aligned} r(s, a, s') &= E[R_t \mid S_{t-1} = s, A_{t-1} = a, S_t = s'] \\ &= \sum_{r \in \mathcal{R}} r p(r \mid s, a, s') \\ &= \sum_{r \in \mathcal{R}} r \frac{p(s', r \mid s, a)}{p(s' \mid s, a, s')} \end{aligned}$$

For episodic tasks, we have nonterminal states \mathcal{S} and terminal states \mathcal{S}^+ . Time of termination T is a random variable between episodes.

Returns at time step t is

$$\begin{aligned} G_t &= R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots \\ &= R_{t+1} + \gamma G_{t+1} \end{aligned}$$

All reinforcement learning algorithm involve estimating 3 quantities: value function, state-action function.

A policy is a probability distribution on \mathcal{A} given s denoted as $\pi(a \mid s)$.

A value function of a state s under a policy π is the expected return when starting in s and following π thereafter.

$$\begin{aligned} v_\pi(s) &= E_\pi[G_t \mid S_t = s] \\ &= E_\pi \left[\sum_{i=0}^{\infty} \gamma^i R_{t+i+1} \mid S_t = s \right], \quad (\text{for all } s \in \mathcal{S}) \\ &= E_\pi \left[R_{t+1} + \gamma \sum_{i=0}^{\infty} \gamma^i R_{(t+1)+i+1} \mid S_t = s \right] \\ &= E_\pi[R_{t+1} \mid S_t = s] + \gamma E_\pi[G_{t+1} \mid S_t = s] \\ &= \sum_{a, r, s'} r p(s', r \mid s, a) \pi(a \mid s), \quad (\text{Law of Total Expectation}) \\ &\quad + \gamma \sum_{a, s'} E_\pi[G_{t+1} \mid S_{t+1} = s'] p(s' \mid s, a) \pi(a \mid s) \\ &= \sum_{a, r, s'} r p(r \mid s, a) \pi(a \mid s) + \gamma \sum_{a, r, s'} v_\pi(s') p(s', r \mid s, a) \pi(a \mid s) \\ &= \sum_a \pi(a \mid s) \sum_{r, s'} p(s', r \mid s, a) [r + \gamma v_\pi(s')] \end{aligned}$$

The state-action function is the expected return of taking an action a at state s before following the policy thereafter.

$$\begin{aligned} q(s, a) &= E_\pi[G_t \mid S_t = s, A_t = a] \\ &= E_\pi \left[\sum_{i=0}^{\infty} \gamma^i R_{t+i+1} \mid S_t = s, A_t = a \right] \\ &= E_\pi[R_{t+1} \mid S_t = s, A_t = a] + \gamma E_\pi[G_{t+1} \mid S_t = s, A_t = a] \\ &= \sum_{s', r} r p(s', r \mid s, a) + \gamma \sum_{s', a'} E_\pi[G_{t+1} \mid S_t = s', A_t = a'] p(s', a' \mid s, a) \\ &= \sum_{s', r} r p(s', r \mid s, a) + \gamma \sum_{s', a'} q(s', a') p(a' \mid s', s, a) p(s' \mid s, a) \\ &= \sum_{s', r} r p(s', r \mid s, a) + \gamma \sum_{s', a', r} q(s', a') \pi(a' \mid s') p(s', r \mid s, a) \\ &= \sum_{s', r} p(s', r \mid s, a) \left[r + \gamma \sum_{a'} q(s', a') \pi(a' \mid s') \right] \end{aligned}$$

Optimal policy and value functions

A policy $\pi' \geq \pi$ iff $v_{\pi'}(s) \geq v_\pi(s), \forall s \in \mathcal{S}$. Optimal state value function is defined as

$$v_*(s) = \max_{\pi} v_\pi(s), \forall s \in \mathcal{S}$$

Similarly, optimal state action value function is defined as

$$q_*(s, a) = \max_{\pi} q_\pi(s, a)$$

Bellman optimality equation for v_*

$$\begin{aligned}
v_*(s) &= \max_a q_{\pi_*}(s, a) = \max_a E_{\pi}[G_t \mid S_t = s, A_t = a] \\
&= \max_a E_{\pi_*}[R_{t+1} + \gamma G_{t+1} \mid S_t = s, A_t = a] \\
&= \max_a [E_{\pi_*}[R_{t+1} \mid s, a] \\
&\quad + \gamma \sum_{s', r} E_{\pi_*}[G_{t+1} \mid S_{t+1} = s'] p(s', r \mid s, a)] \\
&= \max_a \left[\sum_{s', r} r p(s', r \mid s, a) + \gamma \sum_{s', r} v_*(s') p(s', r \mid s, a) \right] \\
&= \max_a \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s')]
\end{aligned}$$

Bellman optimality equation for q_* $q_*(s, a)$ is the value of taking action a at state s , then follow optimal policy after. Hence we can also write.

$$\begin{aligned}
q_*(s, a) &= \max_{\pi} q_{\pi}(s, a) \\
&= \max_{\pi} E_{\pi}[R_{t+1} + \gamma G_{t+1} \mid s, a] \\
&= \max_{\pi} (E_{\pi}[R_{t+1} \mid s, a] + \gamma E_{\pi}[G_{t+1} \mid s, a]) \\
&= \max_{\pi} \left(E_{\pi}[R_{t+1} \mid s, a] + \gamma \sum_{s'} E_{\pi}[G_{t+1} \mid s'] p(s' \mid s, a) \right) \\
&= \left(E_{\pi}[R_{t+1} \mid s, a] + \gamma \max_{\pi} \sum_{s'} v_{\pi}(s') p(s' \mid s, a) \right) \\
&= \sum_{s', r} p(s', r \mid s, a) \left[r + \gamma \max_{\pi} v_{\pi}(s') p(s', r \mid s, a) \right] \\
&= \sum_{s', r} p(s', r \mid s, a) [r + \gamma v_*(s') p(s', r \mid s, a)] \\
&= \sum_{s', r} p(s', r \mid s, a) \left[r + \gamma \max_{a'} q_*(s', a') p(s', r \mid s, a) \right]
\end{aligned}$$

Once we have $v_*(s)$, optimal policy is just a one step search, it selects the action that maximize the sum of immediate reward and the value of the next state i.e. $\max\{R_{t+1} + \gamma v_*(S_{t+1})\}$. If we have $q_*(s, a)$ instead, then for any s , the optimal policy will be just the action that maximize $q_*(s, a)$.

Solving bellman equation is not practical since we all satisfy the following 3 conditions:

1. Know dynamics of the environment (Poker provides partially observed state only)
2. Have enough computation resource (Go has huge state space exceed the number of atoms in the universe)
3. Problem needs to have markov property (Weather prediction cannot be based on yesterday only)

Therefore we typically find approximations to the value functions.

Solving Bellman Equation with DP

Policy Evaluation

Recall that for any policy π ,

$$v_{\pi}(s) = \sum_a \pi(a \mid s) \sum_{r, s'} p(s', r \mid s, a) [r + \gamma v_{\pi}(s')]$$

If the environment is known, we can solve it as a set of linear equations. If not, we can approximate v_{π} given π by initializing v_0 and following the iteration: (Note that unique existence and convergence is guaranteed as long as $\gamma < 1$ (e.g. $\|v_{k+1} - v_k\|_{\infty} \rightarrow 0$).)

$$\begin{aligned}
v_{k+1}(s) &= E_{\pi}[R_{t+1} + \gamma v_k(S_{t+1}) \mid S_t] \\
&= \sum_a \pi(a \mid s) \sum_{r, s'} p(s', r \mid s, a) [r + \gamma v_k(s')]
\end{aligned}$$

Since we are averaging the next states, this is called expected update. There are two ways to update algorithm, "in-place update" where we do a sweep through a single array ($v(s')$ can be the updated value for some s') and "double buffers update" where we keep two arrays, update the back buffer array using the front buffer array and swap the reference. Both will converge. We usually prefer the former.

Policy Improvement